

Lecture 30

Machine Learning: Introduction to Applications of Deep Learning

MP574: Applications

Sean B. Fain, PhD (sfain@wisc.edu)

Diego Hernando, PhD (dhernando@wisc.edu)

ITK/VTK Applications: Andrew Hahn, PhD (adhahn@wisc.edu)

Learning Objectives

- Training and tuning of deep learning networks continued
- Current approaches to analyzing performance
- Some application examples



Tuning Training Process

Batch Learning

- Divides input data into equally sized batches B_k
- The training algorithm processes one batch at a time
- After each batch, the combined cost function is calculated and the weights are updated
- Increases learning rate
 - Smaller batches theoretically allow faster learning
 - If batch size is too small gradient directions might be wrong
 - Iterations might cancel each other out
 - Training might be slow or stagnate

Optimizers

- Gradient descent variations
 - Momentum: combines gradient directions with previous iterations





- AdaGrad: calculates separate learning rate for each parameter (higher learning rates for sparse parameters)
- Adam: combination of Momentum and AdaGrad optimizers

Input Data

- Deep learning requires large amount of training data
- Rule of thumb
 - Training data should be on the order of the number of variables in the model
- More data is always better
- Quality of data matters
 - Has to be representative for future data input
 - Consider possible variables

Generalization and Overfitting

- Just because a trained model works well on the **training data** does not mean it performs well for new data
- Common Problem: Overfitting

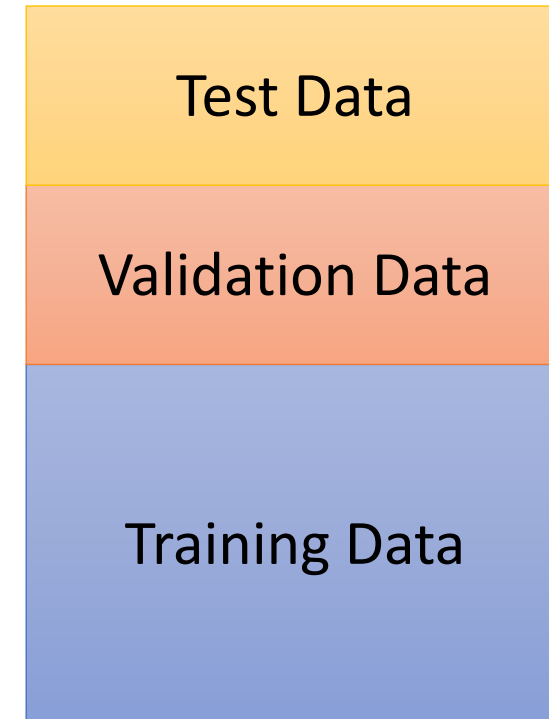


Generalization and Overfitting

- Common causes for overfitting
 - Not enough training data
 - Training data population too narrow
 - Too many variables in the model
 - Too many training iterations
- Monitoring training progress with independent data is important

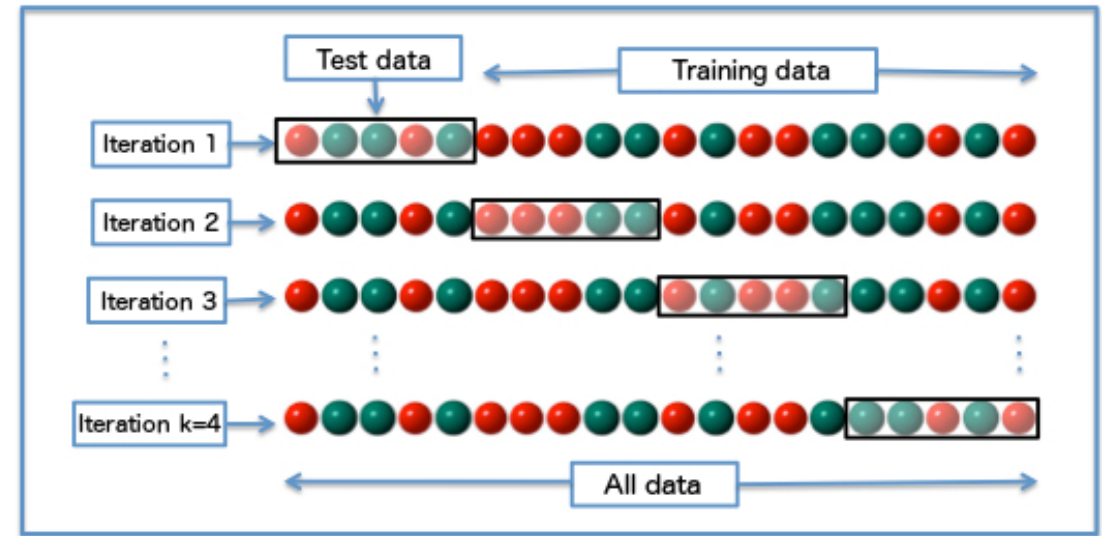
Data (Training, Validation, Test)

- Split data into three groups
 - Training: used to train a neural network (estimate weights)
 - Validation: used to monitor training progress and decide when to stop training process
 - Test: used to test performance and compare different models
- Rule of thumb: 50-25-25 to 60-20-20



Cross-Validation

- Common technique if amount of data is small
- N-fold cross validation
 - Divide data into N groups
 - Perform training with N-1 groups
 - Test trained model with remaining group
 - Repeat process while leaving out a different group for testing
 - Average Results
- Extreme case: leave-one-out




By Fabian Flöck [CC BY-SA 3.0 (<https://creativecommons.org/licenses/by-sa/3.0>)], from Wikimedia Commons

Regularization

- Parameter norm penalties (sparsifying weights)
 - L2-Norm $\|\mathbf{w}\|_2 = \sqrt{\sum_{\forall i} w_i^2}$
 - L1-Norm $\|\mathbf{w}\|_1 = \sum_{\forall i} |w_i|$
 - Add regularization term to cost (“Loss”) function e.g. $L = \|\tilde{\mathbf{y}} - \mathbf{y}\|_2 + \alpha \|\mathbf{w}\|_2$
- Early Stopping
- Bagging: Train multiple models (networks) and vote on output
- Dropout: Randomly exclude non-output nodes for each training sample
 - Approximates bagging





Dataset Augmentation

- Artificially increase amount of data by adding images multiple times with
 - Added noise
 - Geometric transforms (e.g. translation, rotation, scaling)
- Improves robustness of the trained model towards noise and transformations
- Reduces overfitting problem



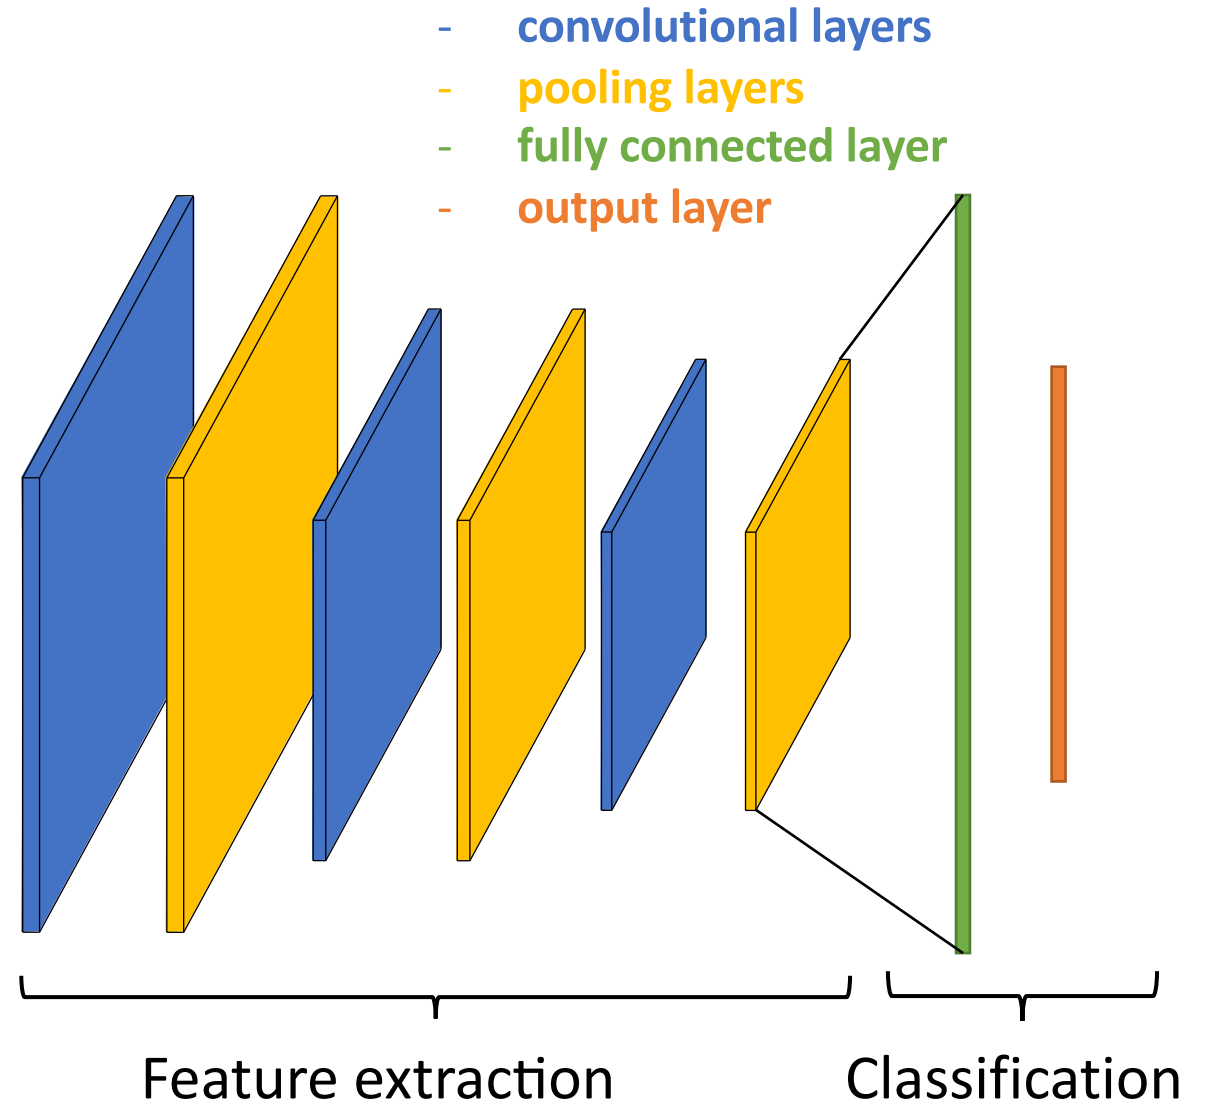
Dataset Augmentation

- Artificially increase amount of data by adding images multiple time with
 - Added noise
 - Geometric transforms (e.g. translation, rotation, scaling)
- Improves robustness of the trained model towards noise and transformations
- Reduces overfitting problem




Transfer Learning

- Idea
 - Features for object recognition are similar for different tasks
 - Lower level layers in CNN determine image features
 - Object recognition / classification is performed in the last layer(s)
 - Reuse pretrained neural network
 - Same architecture and weights for lower level layers
 - Replace output layer by a new layer with appropriate size
 - Initialize weights randomly for the output layer
 - Retrain the model for the new task



Data (Batch) Normalization

- Additional layer
- Normalizes the input of each node of a layer by
 - Subtracting the batch mean
 - Dividing by batch standard deviation
- Goal: speed up learning process
 - Reduce internal covariance shift
 - Allows higher learning rates
 - Potentially decreases overfitting



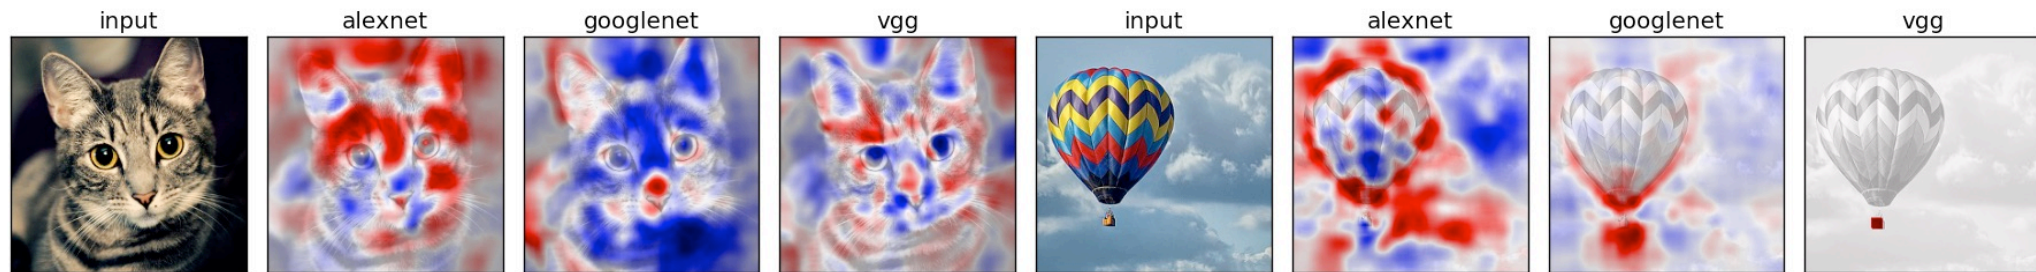
Analyzing Deep Learning Networks

Inside a Convolutional Neural Network

- Common critique: artificial neural networks are **black boxes**
 - Difficult to control what features the network learns
 - May learn inadequate features if input data is not representative of true population
- Understanding how neural networks learn may help
 - Improve the design of networks
 - Improve generalization
 - Make deep learning applications more secure
- Visualizing artificial neural networks is an active field of research

Visualization

- Zintgraf et al. 2017: for a given image detect areas which provide
 - evidence against each class
 - evidence for each class
 - Iteratively removes patches from the input data




L. M. Zintgraf, T. S. Cohen, T. Adel, and M. Welling, “Visualizing Deep Neural Network Decisions: Prediction Difference Analysis,” Nov. 2016.

Visualization


- Zeiler et al.
 - Uses deconvolutional neural network to invert operations
 - Max-pooling operations are not invertible instead algorithm follows location of maximum
 - Visualizes features at each level of the network

Layer 1



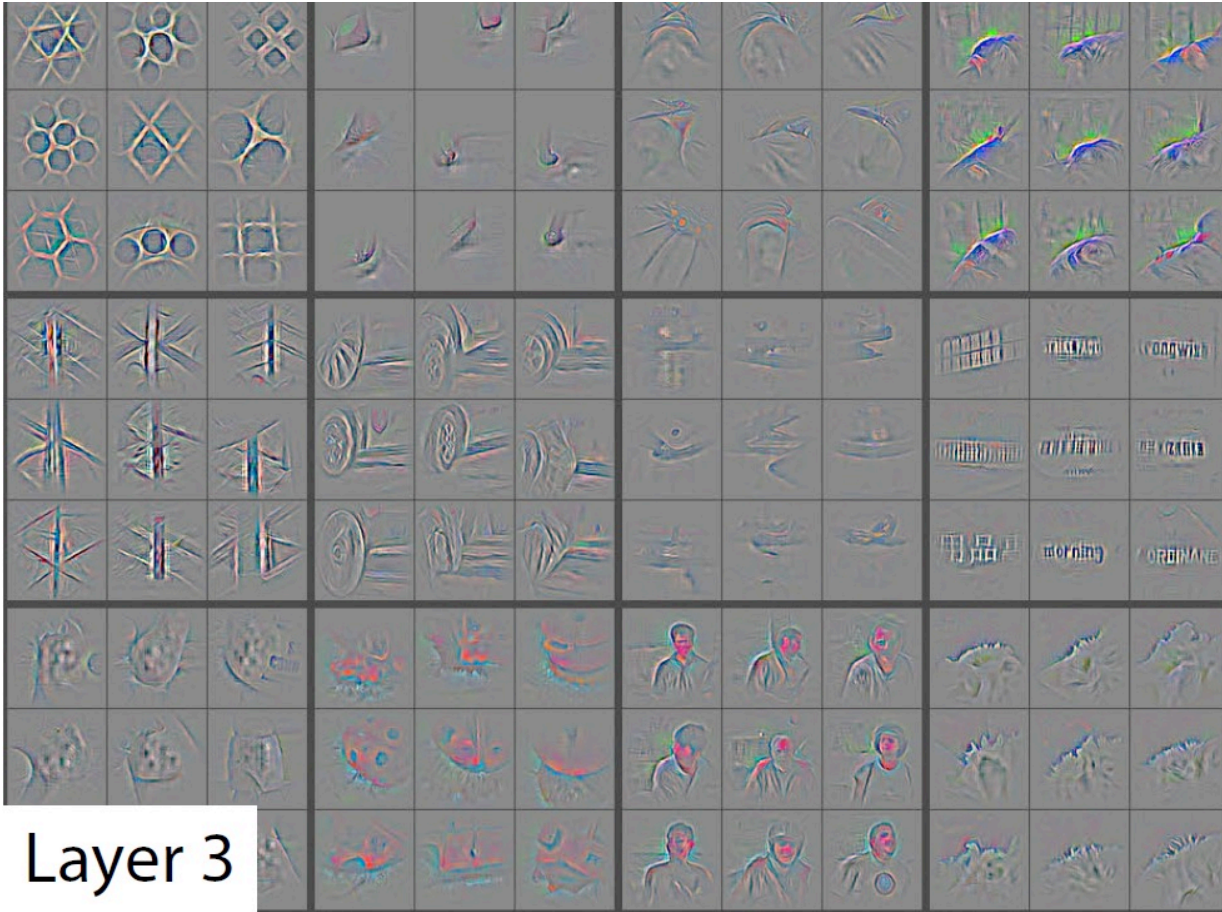
[1] M. D. Zeiler and R. Fergus, “Visualizing and Understanding Convolutional Networks,” in *Computer Vision – ECCV 2014*, vol. 8689, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Cham: Springer International Publishing, 2014, pp. 818–833.

Visualization



[1] M. D. Zeiler and R. Fergus, “Visualizing and Understanding Convolutional Networks,” in *Computer Vision – ECCV 2014*, vol. 8689, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Cham: Springer International Publishing, 2014, pp. 818–833.

Visualization




[1] M. D. Zeiler and R. Fergus, “Visualizing and Understanding Convolutional Networks,” in *Computer Vision – ECCV 2014*, vol. 8689, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Cham: Springer International Publishing, 2014, pp. 818–833.

Segmentation using Deep Learning

Segmentation: Patch Classification

- Classify each pixel of the image separately
- Use local patch around each pixel as input for deep learning network
- Common application:
 - Digital pathology
 - Microscopy
- Only includes local information
 - Not suitable for semantic segmentation of larger complex objects


Example: Neuronal Membrane Segmentation



[1] D. C. Ciresan, L. M. Gambardella, and A. Giusti, “Deep Neural Networks Segment Neuronal Membranes in Electron Microscopy Images,” p. 9.

Cascaded CNN Architecture

- First network creates candidate map
- Second network refines segmentation

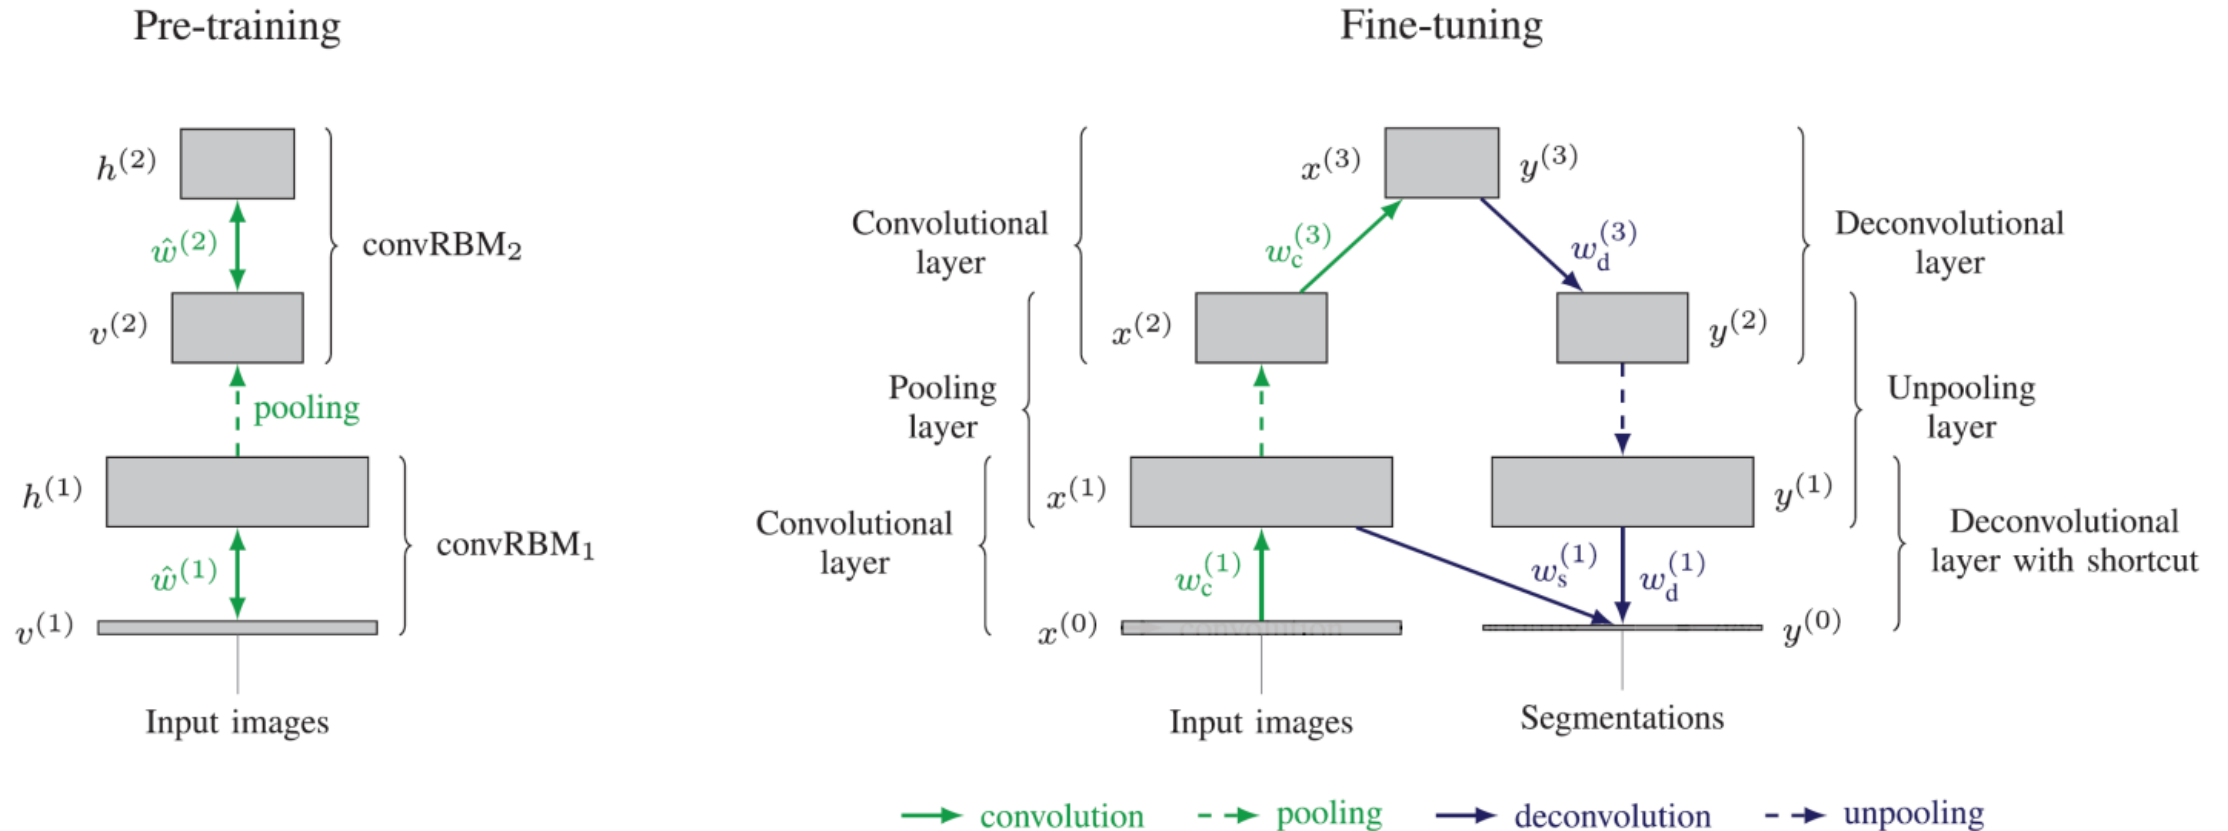


[1] Z. Akkus, A. Galimzianova, A. Hoogi, D. L. Rubin, and B. J. Erickson, “Deep Learning for Brain MRI Segmentation: State of the Art and Future Directions,” *J Digit Imaging*, vol. 30, no. 4, pp. 449–459, Aug. 2017.

Semantic-Wise CNN Architecture

- Deep learning architecture consists of
 - Encoder: convolutional neural network
 - Decoder: deconvolutional neural network
- Output is pixel-wise classification of the input image
- Can learn location dependent information
- Can segment large complex objects

Semantic-Wise CNN Architecture



- [1] T. Brosch, L. Y. W. Tang, Y. Yoo, D. K. B. Li, A. Traboulsee, and R. Tam, "Deep 3D Convolutional Encoder Networks With Shortcuts for Multiscale Feature Integration Applied to Multiple Sclerosis Lesion Segmentation," *IEEE Transactions on Medical Imaging*, vol. 35, no. 5, pp. 1229–1239, May 2016.

Common Preprocessing Steps

- Registration to common anatomical space
- Bone extraction (skull stripping) may help focus segmentation process on soft tissue only
- Bias field correction: correction of image contrast variations due to magnetic field inhomogeneity
- Intensity normalization: similar histogram distributions across different images
- Noise reduction

Summary

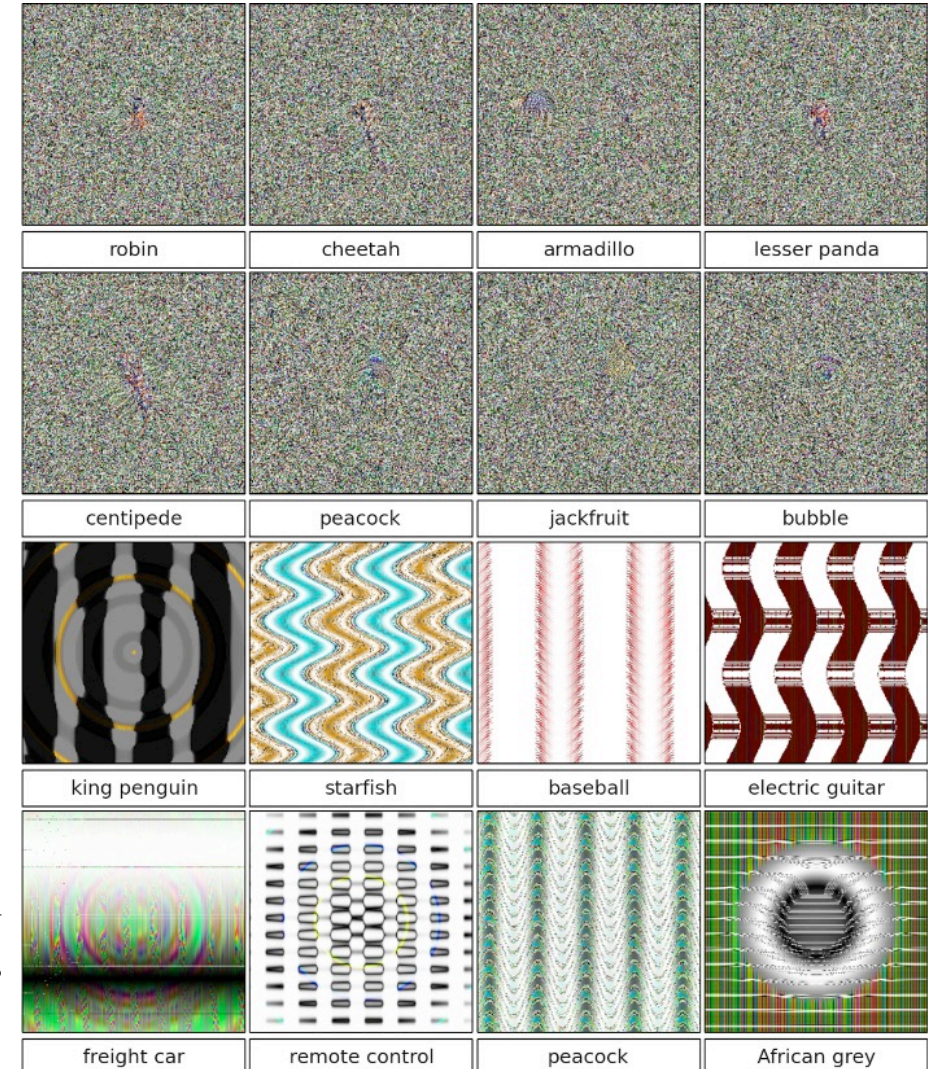
- Deeper understanding of underlying processes of deep learning
 - Network architecture
 - Training algorithm
 - Mathematics behind deep learning
- Many high-level libraries available
 - Understanding the background helps deciding for
 - Network type and architecture
 - Loss functions
 - Activation functions
 - Understanding the importance of input data

Different Types of Deep Learning Networks

Fooling Deep Learning Networks

- Deep convolutional neural networks can often be fooled by textures or noise
- Security risk for
 - Biometric identification (e.g. face recognition)
 - Spam filters
 - Object detection in security cameras

[1] A. Nguyen, J. Yosinski, and J. Clune, “Deep neural networks are easily fooled: High confidence predictions for unrecognizable images,” 2015, pp. 427–436.



Generative adversarial networks

- Introduced by Ian Goodfellow 2014
- Set of two neural networks
 - Generative model: creates artificial input data
 - Discriminative model: trained to discriminate between true and artificial data
- Discriminative model usually pretrained using real data
- Input for generator may be
 - Noise images
 - Text description of the image contents
 - ...

Generative adversarial networks



Salimans, T., Goodfellow, I.J., Zaremba, W., Cheung, V., Radford, A., & Chen, X. (2016). Improved Techniques for Training GANs. *NIPS*.

Text description	This flower has a lot of small purple petals in a dome-like configuration	This flower is pink, white, and yellow in color, and has petals that are striped	This flower has petals that are dark pink with white edges and pink stamen	This flower is white and yellow in color, with petals that are wavy and smooth
64x64 GAN-INT-CLS				
256x256 StackGAN				


H. Zhang, T. Xu, and H. Li, “StackGAN: Text to Photo-Realistic Image Synthesis with Stacked Generative Adversarial Networks,” in *2017 IEEE International Conference on Computer Vision (ICCV)*, 2017, pp. 5908–5916.

Unsupervised Deep Learning

- Labeling images by hand can be time consuming and expensive
- Unsupervised: no labeled images are available for training
- **Theory:** deep neural networks learn by removing irrelevant information in each layer
- Common examples: clustering, estimation of probability distributions
- Unsupervised deep learning:
 - Train a neural network to learn relevant features to describe a set of images


Unsupervised Learning: Autoencoder

- Optimize a network to learn the identity transform
 - Output \hat{x} should be as close as possible to input x
- Hidden layer(s) have less nodes than input and output layers
- Forces network to remove irrelevant information
- Learn only relevant features
- Encoder, decoder scheme similar to image compression



Self-Taught Learning

- Use autoencoder to learn relevant features
- Remove decoder part of autoencoder
- Add new set of output layers
- Train the network with a small set of labeled images
- Similar to transfer learning




Whitening

- Reduce correlation between adjacent blocks of pixels
 - Uncorrelated (orthogonal) basis
 - Equal variance in all directions
- Principle component analysis
 - $\Sigma = \sum_{\forall i} \mathbf{x}_i \mathbf{x}_i^{\top}$
 - Eigenvectors of Σ provide reduced orthogonal basis
- ZCA (Mahalanobis transform)
 - Similar to PCA but rotated to be as close as possible to the original data
- Independent component analysis
 - Maximizes statistical separation between basis vectors

Reinforcement Learning

- Learning through interaction
 - Interacts with environment
 - Observes consequences
 - Alters its behavior in response to rewards received
- Has to learn by trial and error
 - perception-action learning loop
 - Markov decision process
- e.g. learn to play video games directly from pixels




Reinforcement Learning

- Challenges:
 - No ground truth known (only feedback is from rewards)
 - Strong temporal correlations
 - Long range time dependencies
- Value Functions
 - Function that describes the value of a policy given an initial state s
 - The policy which returns the highest value is 'optimal'
 - Often replaced by Quality function which uses initial action
- Policy Search
 - Aims to find policy directly
- Example: AlphaGo
 - Initial training by supervised learning
 - Policy-gradient reinforcement learning
 - Combined using Monte-Carlo tree search


Recursive Neural Networks

- Applies same weights recursively over hierarchical structure
- Used for natural language processing
 - E.g. analyze hierarchical structure of sentences



Recurrent Neural Networks (RNN)

- Special case of recursive neural networks
- Each hidden layer takes
 - Current input
 - Result of previous iteration
- Used for time sequences
 - E.g. natural language processing
- RNNs have a “memory”



By François Deloche [CC BY-SA 4.0
(<https://creativecommons.org/licenses/by-sa/4.0>)], from
Wikimedia Commons

Long Short-Term Memory (LSTM)

- Building block for recurrent neural networks (LSTM unit)
- Composed of
 - Cell state
 - Input gate
 - Output gate
 - Forget gate
- Can remember values over arbitrary long time intervals

